

Alan Huang

(317) 316-7012 | alanhuang@stanford.edu | github.com/alanhuang05 | linkedin.com/in/alanxhuangg

EDUCATION

Stanford University

Stanford, CA

Bachelor of Science in Computer Science

September 2024 – May 2028

- GPA: 3.91 / 4.00

Master of Science in Computer Science (Coterminal)

September 2026 – May 2028

EXPERIENCE

Software Engineer Intern

June 2026 – August 2026

Walmart Global Tech

Bentonville, AR

- Built a Python internal tool for the Walmart Fulfillment Services (WFS) team, used by 40 seller-support agents to look up order status and shipping exceptions, cutting average issue resolution time from 15 to 5 minutes by consolidating queries across 6 backend systems
- Automated ingestion and cleaning of logistics data from 4 warehouse sources using Python (pandas), reducing manual report-prep time from 2 hours to 10 minutes for daily third-party seller performance reports
- Shipped 8 bug fixes and 15 test cases to a production Java-based order-tracking service handling 500,000 orders/day, merged via standard PR review with senior engineers

AI Research Intern

June 2025 – August 2025

National Science Foundation (NSF)

Indianapolis, IN

- Selected for the NSF REU Site: Enhancing Undergraduate Experiences in Data and Mobile Cloud Security
- Improved cross-dataset deepfake detection accuracy from 59.7% to 75.5% on the CelebDF benchmark by designing a diffusion-based Synthetic Domain Alignment technique, validated across FaceForensics++ and CelebDF
- Published as co-first author: *Robust Deepfake Detection via Synthetic Domain Alignment*, 22nd IEEE International Conference on Mobile Ad-Hoc and Smart Systems (IEEE MASS 2025)

PROJECTS

OS: Crash-Consistent Filesystem & Encrypted Virtual Memory | C, C++

2026

- Implemented a Unix V6 filesystem reader in C with write-ahead journaling and crash recovery, restoring consistency after simulated mid-write failures
- Built a memory-mapped encrypted file system in C++ using the Clock page-replacement algorithm to manage 128 physical memory pages, achieving a 27% reduction in page faults compared to a naive FIFO baseline

Multi-Label Hate Crime Bias Classification | Python, scikit-learn, XGBoost

2026

- Trained and benchmarked five classifiers (Logistic Regression, Decision Tree, Random Forest, XGBoost, majority baseline) to predict bias motivation category across 32,508 California hate crime incidents
- Addressed severe class imbalance (three rarest classes <3% of data) using SMOTE and stratified splitting, more than doubling macro F1 from 12.4% (majority-class baseline) to 25.0% with L2-regularized Logistic Regression

Rolyn – Full-Stack Classifieds Platform | JavaScript, Express, SQLite, Supabase

2026

- Built listing detail pages and an authentication backend; implemented secure login using Supabase Auth and JWTs, supporting 50+ user accounts with role-based access across listing and messaging flows.

PUBLICATIONS

Alan Huang, Ian Yee, Shu Hu. *Robust Deepfake Detection via Synthetic Domain Alignment*. The 22nd IEEE International Conference on Mobile Ad-Hoc and Smart Systems (IEEE MASS 2025).

[DOI](#)

Alan Huang, Shaheer Raza, Yuni Xia. *Predicting Diabetes with Medical and Demographic Data*. The 6th International Conference on Computer Science, Engineering, and Education (CSEE 2025).

[DOI](#)

TECHNICAL SKILLS

Languages: Python, C/C++, Java, JavaScript, SQL

ML/Data: TensorFlow, scikit-learn, XGBoost, pandas, NumPy, OpenCV

Systems & Tools: Git, Node.js, FastAPI, Express, Supabase, Google Cloud Platform